



Auditor test–retest stability (5 papers × 10 sealed repeats, frozen code, fixed model/prompt/temperature)


a Finding count is stable




b Reproducible core vs flaky tail




c Reproducibility tracks headlining




d Detection by severity



e Detection by severity × confidence




f Fair-to-moderate (AC1, not α)




merged matcher. n = # distinct issues in each bucket. Severity is a self-label; the high>med step is assignment-rule-sensitive — the robust feature is the low floor.

cell = mean detection rate over 10 runs (merged); n = # distinct issues in the cell (ignore small-n). Severity×confidence is the auditor's top-6 key: 0.74 (high/high) → 0.16 (low/med).

g Flaky ≠ false positive



h Ensembling recovers the misses



i Key findings

What's robust

count stable (CV 0.13–0.22)
low reproducibility = false NEGATIVES,
not false positives (verifier rejects 0)
headlined / blocking defects reproduce
(~0.5–0.7); concrete & verifiable
union 3–5 runs 83–94% recall
reliability fair-to-moderate
(Gwet's AC1 ≈ 0.39)

Do not headline

high>med>low severity gradient
(assignment-rule sensitive; self-label)
Krippendorff α on subsets (paradox)

n=5, repeat-only; merged-by-defect matcher;
GT = verifier + self-checks.